

From Heuristics to Graph Neural Networks: Benchmarking Modularity Maximization for Community Detection

Adil Chhabra
adil.chhabra@informatik.uni-
heidelberg.de
Heidelberg University
Heidelberg, Germany

Kenneth Langedal*
kenneth.langedal@informatik.uni-
heidelberg.de
Heidelberg University
Heidelberg, Germany

Christian Schulz
christian.schulz@informatik.uni-
heidelberg.de
Heidelberg University
Heidelberg, Germany

Abstract

Community detection methods commonly optimize modularity, an objective function designed to maximize intra-community connectivity while penalizing inter-community connectivity relative to random expectation. Over the past two decades, a wide range of algorithms have been proposed to optimize modularity, spanning classical heuristic and exact methods as well as more recent machine learning-based approaches that incorporate node attributes. Yet, survey literature and experimental evaluations remain fragmented, often confined to narrow subsets of algorithms or a single paradigm, limiting insight into relative strengths and trade-offs in practice. In this work, we present the first systematic experimental survey of modularity-based graph clustering algorithms that jointly evaluates classical and learning-based approaches. We study 17 representative algorithms spanning greedy, multilevel, metaheuristic, streaming, exact, and neural paradigms, and organize them within a unified taxonomy grounded in their algorithmic design principles. Through a consistent experimental evaluation across multiple datasets and evaluation criteria, we analyze empirical trade-offs in solution quality and scalability, providing comparative insight rather than identifying a single best method. We additionally contribute a unified, open-source framework for the implementation of all 17 evaluated clustering algorithms, supporting reproducibility and future benchmarking.

CCS Concepts

• **Theory of computation** → **Graph algorithms analysis**; • **Mathematics of computing** → **Graph theory**; • **Computing methodologies** → **Cluster analysis**.

Keywords

Community detection, Graph clustering, Modularity optimization, Experimental evaluation

*Corresponding author.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06

<https://doi.org/XXXXXXX.XXXXXXX>

ACM Reference Format:

Adil Chhabra, Kenneth Langedal, and Christian Schulz. 2018. From Heuristics to Graph Neural Networks: Benchmarking Modularity Maximization for Community Detection. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 10 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

Networks are useful in modeling complex systems across the natural sciences, social sciences, and engineering. Such systems are represented as graphs, where vertices denote entities and edges encode interactions between them, as in social networks, communication infrastructures, biological and financial systems, and information networks. A common feature of real-world graphs is *community structure*, where vertices form groups with dense internal connectivity and sparser external connections.

Identifying these groups, commonly termed *graph clustering* or *community detection*, is a central task in network analysis, with applications ranging from social and political analysis and fraud detection [1] to bioinformatics [13, 50], program optimization [16, 31], and epidemic modeling [34]. Beyond application-specific insights, community structure reveals fundamental organizational patterns, such as cohesive subgraphs, boundary vertices, and inter-community interactions. Despite extensive study, no single clustering method performs best across all network types [37], motivating continued algorithmic development and comparative evaluation.

Most community-detection methods formulate clustering as an optimization problem over vertex partitions, assigning a quality score to each candidate solution. Among clustering objectives, *modularity* [21, 35], which measures the density of edges inside communities as compared to edges between communities while accounting for random chance, is widely adopted and shown to be effective [7, 26]. Thus, this survey focuses on modularity-based methods, while including widely used non-modularity baselines.

We distinguish two broad classes of modularity-based algorithms. *Classical* methods rely on discrete optimization heuristics driven solely by graph topology and a well-defined objective function, yielding interpretable optimization behavior. In contrast, *learning-based* approaches integrate modularity into deep learning pipelines, most commonly via graph neural networks, incorporating node attributes alongside graph topology through feature representations. This added expressiveness introduces trade-offs in interpretability, training complexity, and computational cost [10].

Community detection faces challenges arising from the scale, complexity and frequently-updating content of real-world graphs. Effective methods must balance solution quality, often with noisy

data lacking ground truth, with runtime and memory scalability. Sensitivity to parameters, such as assumed cluster counts, further complicates practical deployment [49]. Understanding how different algorithms navigate these trade-offs remains an open challenge.

Despite existing surveys [4, 24, 27–29, 32, 49], there is no comprehensive experimental study that jointly evaluates classical and learning-based modularity optimization methods under a unified framework. This paper addresses this gap with the following:

Contributions

- We present the first large-scale experimental comparison of classical and learning-based modularity optimization clustering algorithms, with a systematic review of 18 representative methods, spanning greedy, multilevel, metaheuristic, streaming, exact, and learning-based approaches.
- We release a unified open-source experimental framework that adapts both classical and learning-based methods behind a common interface, enabling reproducible benchmarking and future research.
- We propose a taxonomy of modularity-based clustering algorithms, organized by optimization paradigm and algorithmic strategy.

2 Preliminary

Graphs and clusterings. An (*attributed*) *graph* is defined as $G = (V, E, X)$, where $V = \{v_1, v_2, \dots, v_n\}$ is a set of n vertices, $E \subseteq V \times V$ is the set of undirected edges, and $X = \{x_1, x_2, \dots, x_n\}$ denotes node attributes, with each $x_i \in \mathbb{R}^f$ an attribute vector associated with vertex v_i , and f the attribute dimension. We write $n := |V|$ and $m := |E|$ for the number of vertices and edges, respectively. For a vertex $u \in V$, its *neighborhood* is $N(u) := \{v \in V : \{u, v\} \in E\}$, and its *degree* is $\deg(u) := |N(u)|$. Let $A \in \{0, 1\}^{n \times n}$ denote the adjacency matrix of G , where $A_{ij} = 1$ if $\{v_i, v_j\} \in E$ and $A_{ij} = 0$ otherwise. Let $\mathbf{d} \in \mathbb{R}^n$ denote the degree vector with $d_i = \deg(v_i)$.

A *clustering* C of G is a partition of the vertex set V into k disjoint blocks $V_1, \dots, V_k$ such that $V_1 \cup \dots \cup V_k = V$. Classical algorithms operate solely on the graph structure (V, E) , while learning-based methods additionally exploit node attributes X . Learning-based approaches often aim to discover communities that are coherent both in terms of graph connectivity and attribute similarity.

Modularity. The *modularity* score of a clustering is the fraction of edges that are within one cluster minus the expected fraction if the edges were distributed at random. Let $L_i := |\{\{u, v\} \in E : u, v \in V_i\}|$ denote the number of edges internal to block V_i . The modularity of C is defined as

$$Q(C) := \sum_{i=1}^k \left(\frac{L_i}{m} - \left(\sum_{u \in V_i} \frac{\deg(u)}{2m} \right)^2 \right)$$

Equivalently, modularity can be expressed in spectral form. Let $C \in \{0, 1\}^{n \times k}$ denote the (hard) cluster assignment matrix, where $C_{ip} = 1$ if vertex v_i belongs to cluster p and $C_{ip} = 0$ otherwise. Define the *modularity matrix* $B := A - \frac{\mathbf{d}\mathbf{d}^\top}{2m}$. Then modularity can be written as

$$Q(C) = \frac{1}{2m} \text{Tr}(C^\top B C).$$

Table 1: Comparison of benchmarked methods. Code indicates implementation availability (×: unavailable; *: available but not ready for use; **: open source with well-defined I/O). Features indicate parallelism (P), automatic or user-specified number of clusters (A/K), learning-based methods (L), use of node features (F), and availability in portable Python packages (Py). The bold methods are exact solvers.

Method	Year	Code	Features
CNM [15]	2004	**	. A . . Py
WALKTRAP [38]	2005	**	. A . . Py
WT [47]	2007	**	. A . . Py
LOUVAIN [11]	2008	**	. A . . Py
CGGCI [36]	2012	×	. A . . .
VNS [2]	2012	×	. A . . .
COMBO [43]	2014	**	. A . . Py
ADVNDs [18]	2017	×	. A . . .
HOLLOCOU [23]	2017	**	. A . . .
VIECLUS [9]	2018	**	P A . . .
LEIDEN [44]	2019	**	. A . . Py
DMON [45]	2020	*	P K L F .
COMM DGI [52]	2020	*	P K L F .
GNNs [42]	2022	*	P A L . .
UCoDe [33]	2023	*	P K L F .
DGCLUSTER [8]	2023	*	P A L F .
MAGI [30]	2024	*	P K L F .
BAYAN [5]	2024	**	. A . . .
CLUStRE [14]	2025	**	. A . . .
LIM [48]	2025	*	P A L F .
MaxSAT [3]	2025	×	. A . . .

This spectral formulation serves as the foundation for many continuous and learning-based approaches.

Graph neural networks. Graph neural networks (GNNs) are a class of neural architectures that compute vector representations for vertices in an attributed graph by iteratively aggregating information from local neighborhoods. In the context of community detection, GNNs are used to map vertices into a representation space in which structurally or semantically related vertices are close. Learning-based clustering methods then infer communities indirectly by optimizing objectives such as continuous or differentiable relaxations of modularity, mutual-information maximization, or contrastive criteria defined over these representations or over derived *soft cluster assignments*, which encode the strength of association between nodes and communities.

3 Background and Methods

We review objective functions and algorithmic paradigms for community detection and introduce the taxonomy used in this work.

Many clustering quality measures have been proposed, including *coverage*, *performance* [20], *inter-cluster conductance* [25], *surprise* [6], and *modularity* [21]. Coverage, the fraction of intra-cluster edges, is simple but favors trivial solutions. Modularity mitigates this by comparing observed coverage to its expectation under a




Figure 1: Previous experimental comparisons visualized. Each directed edge indicates that the source algorithm used the destination algorithm in the experimental evaluation. The colors indicate either a learning-based method when red, or a classical optimization algorithm when blue.

null model, and has become one of the most widely used objectives for community detection [26]. It is the primary quality function in large-scale benchmarks such as the DIMACS challenges [7] and is the focus of this work. While Aref et al. [4] report that modularity heuristics rarely attain optimal solutions, Ans’otegui et al. [3] show that the VIECLUS’s modularity maximization heuristic routinely computes optimal solutions on small to medium-sized graphs. Alternative objectives capture complementary notions of community structure: *Weighted Community Cluster* (WCC) maximizes the number of triangles within each community [39], implemented in the SCALABLE COMMUNITY DETECTION algorithm [40], and the *Map Equation*, which measures the expected description length of a random walk [41], used in INFOMAP [41], INFOFLOW [19], and the recent learning-based NEUROMAP [10]. We include publicly available methods based on these objectives as non-modularity baselines.

Classical modularity-based algorithms use discrete optimization heuristics that consider graph topology only. More recently, learning-based approaches to community detection have been proposed, motivated by the development of GNNs that can model node attributes as feature vectors to capture additional graph information beyond topology alone. These methods typically optimize training objectives that are explicitly aligned with modularity or related structural principles. Compared to classical heuristics, learning-based methods trade exact modularity gains for flexibility, the ability to incorporate node attributes, and increased generalization across graphs. Despite extensive work in both paradigms, existing surveys do not provide a unified experimental comparison of classical and learning-based modularity optimization methods. Prior surveys typically focus on either classical heuristics [4] or learning-based methods [24, 28, 29, 32]. Li et al. [27] compare classical and deep learning methods, but limit comparison to Python-based solvers, thus lacking a comprehensive comparison of the state-of-the-art. In addition, algorithmic papers in both paradigms rarely evaluate against methods outside their own class (see Fig. 1).

This work addresses this gap by providing a comprehensive experimental comparison of modularity-based community detection algorithms under a unified framework. Here, we provide a taxonomy of the considered methods listed in Table 1. We organize methods by *optimization paradigm*, distinguishing between *classical* combinatorial algorithms and *learning-based* approaches with trainable parameters. Classical methods are further grouped by their

dominant optimization mechanism (e.g., greedy, multilevel, or meta-heuristic) and by computational setting (e.g., streaming and exact solvers). Learning-based methods are categorized by how modularity is used during training, separating approaches that directly optimize continuous relaxations of modularity from those that use modularity as a guiding signal within broader learning objectives. Within each subgroup, algorithms are listed chronologically.

3.1 Classical Modularity Optimization Methods

3.1.1 Greedy agglomeration and multilevel local search.

Early approaches to modularity maximization focused on greedy agglomeration, merging communities to directly increase modularity. Subsequent work identified limitations of purely agglomerative strategies, notably unbalanced merges and poor scalability, shifting toward multilevel formulations based on local vertex moves.

Clauset–Newman–Moore (CNM) [15]. CNM is a hierarchical agglomerative algorithm that greedily maximizes modularity. Starting from singleton communities (each vertex assigned to its own cluster), it repeatedly merges the pair of communities with the largest modularity gain. The algorithm produces a dendrogram representing the full merge hierarchy and returns the clustering attaining maximum modularity. Its running time is $O(md \log n)$, where d is the dendrogram depth.

Wakita–Tsurumi (WT) [47]. WT refines CNM by addressing its tendency to produce highly unbalanced merges that degrade scalability on large graphs. While preserving the same agglomerative structure, WT modifies merge selection to account for both modularity gain and a community-size ratio, biasing the process toward more balanced agglomerations and improving scalability.

LOUVAIN [11]. LOUVAIN replaces community-pair merging with greedy local vertex moves. Vertices are iteratively moved to neighboring communities whenever the induced modularity gain is positive. Once no improving move exists, communities are contracted into supernodes and the process is repeated, yielding a multilevel optimization scheme. Each local-moving phase runs in time linear in the number of edges, and the number of multilevel passes is typically small, resulting in near-linear runtime in practice.

LEIDEN [44]. LEIDEN addresses LOUVAIN’s tendency to produce disconnected communities by introducing a refinement phase before contraction. This phase splits and restructures communities

through probabilistic, non-greedy merges, while guaranteeing that resulting communities are connected and locally optimal.

3.1.2 Non-modularity baselines. We include two widely used non-modularity baselines which rely on random walks to model how information or flow propagates through the network, providing an alternative perspective to modularity optimization.

WALKTRAP [38]. WALKTRAP performs hierarchical clustering based on the intuition that random walks on a graph tend to get "trapped" into densely connected regions corresponding to communities. Each vertex is characterized by the probability distribution over vertices reached by a random walk starting from it, and a structural similarity distance between vertices (and communities) is defined by comparing these distributions. Communities are then obtained by agglomerative merging of the closest pairs under this distance, producing a dendrogram similar to CNM. The algorithm runs in $O(mnd)$ time, where d is the depth of the resulting dendrogram.

INFOMAP [41]. Unlike prior approaches, INFOMAP formulates community detection as an information-theoretic problem, seeking a clustering that minimizes the expected description length of a random walk under the map equation. It uses a greedy multilevel search inspired by LOUVAIN, augmented with additional fine- and coarse-tuning steps to escape local minima, similar to LEIDEN.

3.1.3 Metaheuristics and ensemble/hybrid search. Beyond greedy and multilevel heuristics, several methods apply metaheuristic or ensemble-learning strategies to modularity maximization. Representative examples include Variable Neighborhood Search (VNS) [2], which systematically perturbs the current clustering to escape local optima, and ensemble-based approaches such as CGGCi-RG [36], which stabilize modularity optimization by consolidating recurring community structures across multiple runs. These methods combine diversification mechanisms (e.g., randomized perturbations or multiple initializations) with intensification via modularity-based local refinement. While more computationally expensive than purely greedy algorithms, such strategies often achieve higher modularity by exploring a larger search space.

COMBO [43]. Instead of relocating individual vertices as in LOUVAIN, COMBO considers pairs of communities and evaluates full redistributions of vertices between them using Kernighan–Lin–style refinements. By unifying merges, splits, and partial recombinations, COMBO can escape local optima that trap multilevel greedy methods, at the cost of higher computational complexity. The algorithm runs in $O(n^2C)$ time, where C is the number of clusters.

VIECLUS [9]. VIECLUS further extends ensemble-based optimization by adopting a memetic (evolutionary) framework. Rather than refining a single partition, it maintains a population of clusterings and generates new solutions through recombination and mutation. Recombination exploits agreement between high-quality solutions via overlay clustering, where two vertices are grouped together only if they co-occur in both parent solutions, while mutation introduces controlled perturbations to preserve diversity. Each offspring is refined using LOUVAIN or multilevel local search and retained only if it improves modularity. After a set time limit, VIECLUS returns the best clustering found thus far.

3.1.4 Streaming algorithms. Most classical modularity algorithms assume random access to the full graph, which becomes infeasible for large or dynamically generated networks. Streaming algorithms address this limitation by processing the graph incrementally under strict memory constraints, typically allowing only a single or few passes over the data.

HOLLOCOU [23]. HOLLOCOU is a one-pass edge-streaming algorithm based on the observation that, under random stream orderings, intra-community edges tend to arrive earlier than inter-community edges. The algorithm starts from the singleton clustering and processes edges sequentially. For an arriving edge (u, v) , HOLLOCOU decides whether to merge the communities of u and v based on their current community volumes. If both volumes are below a fixed threshold, the node belonging to the smaller-volume community joins the other. The authors show that this local, volume-based merging rule corresponds to a greedy optimization of modularity in the streaming setting.

CLUStRE [14]. CLUStRE follows the node-streaming model unlike the edge-streaming model of HOLLOCOU. Here, nodes arrive sequentially together with their neighborhood lists. Starting from singleton clustering, each node is assigned to the cluster among its neighbors that maximizes the modularity gain similar to a LOUVAIN local move. During streaming, a quotient graph is constructed whose nodes represent clusters and whose weighted edges summarize inter- and intra-cluster connectivity, preserving modularity exactly. This quotient graph is optionally refined using VIECLUS. The graph may be re-streamed to perform further local improvements, iterating until modularity converges.

3.1.5 Exact solvers. In contrast to heuristic approaches, exact solvers provably compute globally optimal modularity clusters. Since modularity maximization is NP-hard [12], these methods rely on exponential-time search augmented with bounding techniques or logical formulations such as weighted partial MaxSAT [3], and are therefore limited to small or medium-sized graphs.

BAYAN [5]. BAYAN is an exact branch-and-cut algorithm for modularity maximization. It formulates modularity as an integer program with binary co-membership variables and enforces consistency via triangular constraints on node triples, which branch on whether three nodes must belong to the same community or cannot all be clustered together. Upper bounds are obtained by solving linear programming relaxations, while lower bounds are computed from feasible clusterings produced by COMBO. Branching decisions either contract node triples into supernodes or discourage their co-assignment, and the algorithm terminates once optimality is certified or a prescribed optimality gap is reached.

3.2 Learning-Based Methods

3.2.1 Continuous and differentiable modularity maximization. These approaches replace discrete community assignments with continuous representations, often learned using GNNs, expressing modularity as a continuous or differentiable objective optimized over node representations or soft cluster assignments.

DMoN [45]. DMoN introduces a differentiable relaxation of modularity and realizes it using a GNN. A GNN followed by a softmax

function produces soft cluster assignments, which are optimized using an objective whose main term corresponds to modularity in spectral form. The objective further includes a *collapse regularization* term that prevents trivial solutions in which all nodes are assigned to the same cluster.

GNNS [42]. GNNS formulates modularity maximization as a continuous optimization process over soft cluster assignments. Starting from random assignments, node assignments are iteratively updated by combining previous assignments with modularity-weighted aggregation of neighboring assignments. Activation functions and normalization steps amplify valid assignments supported by positive modularity contributions. The procedure is repeated for multiple random initializations and parameter choices, and the best converged solution is discretized to obtain the final clustering.

DGCLUSTER [8]. Most learning-based modularity methods, including DMO, require the number of communities to be specified a priori (see Table 1). DGCLUSTER eliminates this requirement by learning node embeddings with a GNN and clustering implicitly through a joint optimization objective. Node embeddings learned by a GNN induce soft pairwise community assignments, which parameterize a differentiable modularity objective. A secondary loss incorporates auxiliary local information, such as partial labels or clusterings from classical heuristics, by aligning embedding similarities with a pairwise similarity matrix. After end-to-end optimization, final communities are obtained by clustering the learned embeddings using BIRCH [51].

LIM (Less Is More) [48]. While DGCLUSTER removes the need to predefine the number of communities, it relies on jointly optimizing multiple objectives which increases model and training complexity. LIM addresses this trade-off with a single-objective formulation that still does not require prior specification of the number of communities. The method first applies LOUVAIN to extract coarse global communities, which are summarized as structural community centers in a feature space. A graph convolutional network then learns node representations that capture local structural information and node attributes. Soft cluster assignments are computed by measuring similarity between node representations and the structural centers, and a single modularity objective defined over these assignments is optimized to obtain the final clustering.

3.2.2 Contrastive/infomax objectives. These methods do not optimize modularity directly, but instead use it to guide representation learning through alternative objectives.

COMM DGI (Community Deep Graph Infomax) [52]. COMM DGI learns community-aware node representations by jointly maximizing mutual information between nodes and the graph and between nodes and their communities. Unlike other methods that directly optimize modularity, modularity is used here as an auxiliary objective together with mutual information, via a soft k -means-style clustering layer, to encourage structurally coherent communities.

UCoDe (Unified Community Detection) [33]. UCoDe optimizes a single contrastive objective derived from the modularity matrix, encouraging high similarity within communities and low similarity across communities. This contrastive formulation allows UCoDe to recover non-overlapping communities while naturally

supporting overlap, without requiring separate objectives or post-processing. Compared to combined-objective methods such as COMM DGI, UCoDe uses a single, interpretable loss and requires minimal hyperparameter tuning.

MAGI [30]. MAGI establishes a formal connection between modularity maximization and graph contrastive learning by showing that modularity naturally defines positive and negative node pairs via the sign of the modularity matrix. Specifically, node pairs with positive modularity coefficients are treated as positive samples, while pairs with non-positive coefficients are treated as negatives, yielding a contrastive objective equivalent to relaxed modularity maximization. A GNN encoder is trained using a community-aware contrastive loss, with mini-batch modularity matrices constructed via random-walk sampling for scalability.

4 Evaluation Methodology

In this section, we describe the benchmarking methodology used to evaluate modularity-based graph clustering algorithms under controlled and reproducible conditions. A central contribution is the integration of 17 clustering algorithms behind a unified benchmarking interface, enabling a fair comparison across diverse implementations. In addition to enabling the experiments presented in this survey, it will also serve as an openly available resource for future work on graph clustering, simplifying the evaluation of older methods on new data or machines. This section also includes details on the datasets we used to evaluate the clustering algorithms and the computational environment.

4.1 Benchmark Interface

As shown in Table 1, over half of the clustering algorithms in the literature lack ready-to-use interfaces for clustering new data. Furthermore, among those with properly documented code, there are wide differences in the input graph format, output formats, and the additional information they report, such as time and memory consumption. This lack of standardization complicates fair benchmarking, as differences in preprocessing, reporting, and resource accounting can substantially affect observed performance.

To address this problem, we aim to provide an interface with a single common input and output format and consistent time and memory monitoring. The interface exposes a common set of inputs, while allowing solver-specific parameters to be provided when applicable.

$$\text{CLUSTER}(G, T_{\max}, p, k, M_{\max}, \mathcal{F})$$

Where G is the input graph, $T_{\max}$ is the wall-clock time limit, p is the number of threads available to the algorithm, k is the number of clusters to compute, $M_{\max}$ is the memory limit in gigabytes, and $\mathcal{F}$ are the graph features. Parameters p , k , and $\mathcal{F}$ are only used by a subset of the solvers. From Table 1, parallel solvers use p , those that compute a fixed number of clusters use k , and any solver that uses graph features can optionally accept $\mathcal{F}$. When $\mathcal{F}$ is not provided, but the solver expects features, a single 1.0 value is assigned to each vertex as the feature. This interface is designed to capture the most common operational requirements across clustering algorithms while enforcing consistent resource limits and output reporting.

Regarding code changes, we have tried to make as few modifications to the original source code as possible. Whenever possible, we use the authors' implementations or the most commonly used library implementation. The exact setup for each algorithm is described in our openly available GitHub repository¹. We make only minimal, deliberate modifications where necessary to ensure consistent timing, memory accounting, and output quality measures based on the following rules:

- Parsing the input and writing the output to a file is not counted in the timing of the algorithm.
- For search-based algorithms that always run for the entire $T_{\max}$ duration, the recorded computation time is defined as the point in time when the best solution was found.
- Learning-based methods that train a neural network as part of the algorithm will break the training loop at $T_{\max}$, and use the current model at that time to produce the final clustering.

4.2 Parameter Settings

As much as possible, we try to use the default configurations for all algorithms tested. In the following we detail the exact parameters chosen for each algorithm. Algorithms that are not listed do not expose any parameters.

Parameter Selection Policy. To reflect how algorithms are typically used in practice, we deliberately avoid any dataset-specific or instance-specific parameter tuning. Our goal is to benchmark available implementations as they are, rather than hypothetical optimal configurations that would require extensive tuning effort.

- WALKTRAP: walk length = 4.
- LOUVAIN: resolution = 1.
- HOLLOCOU: $v_{\max} = 10\,000$.
- VIECLUS: time limit = 1 hour.
- LEIDEN: resolution = 1, $\beta = 0.02$, iterations = 2.
- DMOON: 1 hidden layer (64), 1000 epochs, lr = 0.001.
- COMMDOG: 1 hidden layer (512), 1000 epochs, lr = 0.001.
- GNNS: GNNS100 variant.
- UCODE: 2 hidden layers (256), 1000 epochs, lr = 0.001.
- DGCLUSTER: hidden layers (256, 128), 300 epochs, lr = 0.001.
- MAGI: hidden layers (512, 256), 100 epochs, lr = 0.01.
- BAYAN: Gurobi ILP solver [22], gap = 10^{-4} , resolution = 1.
- CLUSTRE: light mode, modularity objective.
- LIM: 1 hidden layer (512), 1000 epochs, patience = 200, lr = 0.001.

For stochastic algorithms with an explicitly specified seed, we use the run ID as the seed for each of the five runs. While many of these algorithms include multiple configurations, we use either the configuration recommended by the authors or the default configuration provided with the implementation.

4.3 Dataset

We consider three datasets commonly used in the literature. The first set consists of 10 real-world graphs with known community labels; we refer to this dataset as the GROUND-TRUTH graphs. The machine learning algorithms commonly use these graphs, and the

¹<https://github.com/KennethLangedal/Clustering>

Table 2: Real-world instances with known ground-truth clusters. The $|C|$ columns indicates the number of clusters.

Instance	n	m	$ \mathcal{F} $	$ C $
Amazon-Computers	13 752	245 861	767	10
Amazon-Photo	7 650	119 081	745	8
Citeseer	3 327	4 552	3 703	6
Coauthor-CS	18 333	81 894	6 805	15
Coauthor-Physics	34 493	247 962	8 415	5
Cora	2 708	5 278	1 433	7
ogbn-Arxiv	169 343	1 157 799	128	40
ogbn-Products	2 449 029	61 859 012	100	47
Pubmed	19 717	44 324	500	3
Reddit	232 965	57 307 946	602	41

details of the instances are shown in Table 2. Next, we consider the instances used in the 10th DIMACS implementation challenge on graph clustering; we call this dataset DIMACS graphs. This dataset is popular among the classical modularity optimization methods, and the details of these graphs are shown in Table 3 in the appendix. Finally, we use synthetically generated graphs with node features generated using the attributed, degree-corrected stochastic block model [46] (ADC-SBM). This generator plants k clusters in the graph, and generates random edges via a distribution that is conditional on those clusters. It does so using the two input parameters d and $d_{out} \leq d$, where d is the expected degree for each vertex and d_{out} is the expected number of neighbors for a node in a different cluster than its own. The ADC-SBM generator also produces node features that are correlated with the community structure. It does so using the two input parameters σ for intra-cluster center variance and σ_c for cluster center variance.

4.4 Computational Environment

All algorithms are evaluated on the same machine to eliminate variability due to hardware differences and ensure fair comparison. While some methods could potentially benefit from specialized accelerators such as GPUs, we intentionally restrict all experiments to a shared CPU-based environment to maintain consistency across algorithms. We use a machine with an AMD EPYC 9754 128-core processor with 755 GB of memory, running Ubuntu 24.04.3 with Linux kernel 6.8.0-87. For algorithms with compiled software, we use GCC version 13.3.0, and for Python implementations we use Python version 3.12.3.

In order to speed up experiments, we evaluate three algorithms in parallel. For parallel algorithms, we evaluate them both sequentially and using 16 threads. If an algorithm exceeds 250 GB of memory usage on an input instance, it is killed and recorded as *memory limit exceeded* for that graph. Similarly, we assign each algorithm a time limit of 1 hour per input graph. Unlike the memory limit, we allow limited flexibility for time limits due to the structure of learning-based methods, which often train for a fixed number of iterations. Only if an algorithm overshoots the time limit by more than 50% do we kill the process and record that instance as *time limit exceeded*.

For algorithms that need to be told the number of clusters, we either provide the exact number when ground truth labels are known, or 16 clusters otherwise, as is commonly used in the literature. Since several of the algorithms are Stochastic, we report the best result over five runs to reflect the maximum achievable solution quality under the given resource constraints, which is standard practice in modularity optimization. In addition to measuring modularity as defined in Section 2, we also use Normalized Mutual Information (NMI) between the computed clusters and ground-truth labels and F1 scores between adjacent pairs of vertices.

Performance Profiles. Performance profiles are well suited for benchmarking multiple algorithms, as they summarize relative performance across instances without requiring direct aggregation of objective values [17]. A performance profile assigns each included algorithm a non-decreasing, piecewise-constant function. The y-axis represents the fraction of graphs where the objective value is better than or equal to τ times the best observed objective function value from any algorithm. Reading a performance profile, an algorithm performs well if it achieves lower τ values for higher fractions of instances.

5 Experiments

In this section, we present benchmark results obtained using our unified benchmarking interface described in Section 4. We report results on real-world graphs with and without ground-truth labels, as well as synthetically generated graphs with controlled noise in structure and features. Rather than exhaustively ranking algorithms, these experiments aim to illustrate how different algorithmic design choices behave under a standardized benchmarking setup.

5.1 State-of-the-Art Comparisons

Figure 1 shows the main results of our experimental comparison. Across both datasets, VIECLUS achieves the highest modularity whenever it completes within the imposed resource limits, consistent with prior findings [3, 9]. As expected for a memetic search-based method, it exceeds the memory limit on the largest DIMACS instances.

On the GROUND-TRUTH dataset, many methods achieve similar high F1 scores. Among them, MAGI achieves the strongest and most stable performance. This highlights a general mismatch between modularity optimization and externally defined communities. For example, the Citeseer ground-truth clustering has a modularity of 0.54, whereas all methods produce clusterings with modularity above 0.54. This discrepancy reflects the fact that citation behavior does not necessarily align with categorical labels, as Figure 1 also indicates.

For the DIMACS graphs, differences between methods are more pronounced, with modularity-focused algorithms clearly outperforming learning-based approaches. Among learning-based methods, performance varies substantially depending on whether the method is designed to operate without node features, with GNNS performing best in this setting.

The streaming methods HOLLOCOU and CLUSTRE trade solution quality for extreme efficiency. While their F1 scores are lower on the GROUND-TRUTH dataset, CLUSTRE achieves competitive modularity on DIMACS instances while using orders of magnitude less

time and memory. The exact algorithm BAYAN terminates only on the eight smallest instances from the DIMACS dataset and is therefore excluded from the performance profiles. For these graphs, VIECLUS found the optimal clustering for all of them, LEIDEN for 7/8, GNNS for 5/8, COMBO for 4/8, and LOUVAIN for 3/8. The remaining algorithms did not find any known optimal solutions.

5.2 Quality-Efficiency Tradeoffs

Figure 3 summarizes the tradeoff between clustering quality and execution time using geometric means across instances. Streaming methods occupy the extreme low-latency regime, while greedy modularity methods (LOUVAIN, LEIDEN) achieve a favorable balance between quality and efficiency. The search-based algorithm VIECLUS achieves the highest quality, but at a substantially higher computational cost. Similarly, learning-based methods generally attain high solution qualities at substantially higher computational cost.

Peak memory usage varies by several orders of magnitude across methods; we provide detailed measurements for the GROUND-TRUTH dataset in Appendix B. We also evaluate the parallel algorithms sequentially to show how additional compute resources affect these methods; we provide detailed measurements in Appendix C.

5.3 Synthetic Graphs

To further evaluate the impact and sensitivity of graph features across clustering algorithms, we include experiments with synthetically generated graphs with vertex features. These experiments illustrate how feature-aware methods respond to controlled degradation of structural and attribute signals, rather than ranking the methods. For reference, we also include VIECLUS, which performed best among the pure modularity-optimizing methods. We conduct two sets of experiments with the ADC-SBM generator. First, we generate a sequence of graphs with gradually higher d_{out} , causing the graph structure to carry less information about the original k clusters. And second, we generate graphs with strong community structure as we gradually increase the cluster center variance σ_c , leading the graph features to carry less information about the clusters.

Figure 4 summarizes the results of the two ADC-SBM experiments. When the feature signal is strong, MAGI and DMOON achieve the highest NMI scores, and beyond $d_{out} \approx 11$, all feature-aware methods outperform VIECLUS. The results also highlight clear differences in how methods balance structural and feature information. MAGI relies primarily on node features, as evidenced by its strong performance under structural noise and its sensitivity to feature degradation. In contrast, DMOON remains robust in both settings, indicating a more balanced use of topology and attributes. COMDGI exhibits similar trends but performs consistently worse than DMOON. DGCLUSTER shows the opposite bias to MAGI, favoring graph structure over features, while UCODE displays a similar but less pronounced tendency. We exclude LIM from this figure, as it consistently collapsed to a single cluster, yielding zero NMI across all configurations.




Figure 2: Performance profiles for our first two datasets. We use F1 scores for the GROUND-TRUTH set and modularity for the DIMACS instances. For clarity, we split the x-axis to highlight the differences between the best algorithms near one.




Figure 3: Scatter plot showing the ratio between running time and solution quality, using the geometric mean for both measures. To normalize the running time for each graph, we divide it by the number of edges, yielding a time-per-edge measure.

6 Conclusion

We surveyed over twenty modularity-based graph clustering algorithms, covering classical heuristics, parallel methods, and recent learning-based approaches. When modularity maximization is the sole objective, the results are clear: **VIECLUS** consistently achieves the highest modularity values when sufficient memory is available. **LEIDEN**, an improved variant of **LOUVAIN**, closely follows in solution quality while offering substantially faster runtimes. This trade-off reflects their optimization strategies: **VIECLUS** relies on time-budgeted search, whereas **LEIDEN** employs a greedy refinement scheme.

Beyond modularity alone, other algorithms exhibit their own unique strengths. Streaming methods such as **CLUSTRE** offer extremely low runtimes and memory footprints, making them well-suited for large-scale or resource-constrained settings. Learning-based approaches, including **GNNS**, demonstrate that neural methods are approaching the performance of hand-crafted heuristics. However, they still require substantially more memory and do not yet achieve competitive modularity scores. At the same time, there is growing interest in expanding modularity optimization with auxiliary objectives to recover ground-truth communities better. In this setting, feature-aware methods can be advantageous: **MAGI** achieves the strongest ground-truth recovery in our experiments, while **DMoN** remains effective when community structure is encoded between topology and node attributes.

To support future work in this area, we introduce a unified benchmarking framework that enables direct, reproducible comparisons across a wide range of modularity-based clustering algorithms. We hope that this framework lowers the barrier for evaluating new methods against established baselines and encourages more systematic experimental studies of graph clustering algorithms.

Acknowledgments

We acknowledge support by DFG grant SCHU 2567/5-1.

References

- [1] Leman Akoglu, Hanghang Tong, and Danai Koutra. 2015. Graph based anomaly detection and description: a survey. *Data Min. Knowl. Discov.* 29, 3 (2015), 626–688. doi:10.1007/S10618-014-0365-Y
- [2] Daniel Aloise, Gilles Caporossi, Pierre Hansen, Leo Liberti, Sylvain Perron, and Manuel Ruiz. 2012. Modularity maximization in networks by variable neighborhood search. In *Graph Partitioning and Graph Clustering*. Contemporary Mathematics, Vol. 588. American Mathematical Society, 113–127. doi:10.1090/conm/588/11705
- [3] Carlos Ansótegui, Vaidyanathan Peruvemba Ramaswamy, Stefan Szeider, and Hai Xia. 2025. Uncovering and Verifying Optimal Community Structure in Complex Networks: A MaxSAT Approach. In *Proceedings of the 25th International Conference on Computational Science (ICCS 2025)*. Springer Nature Switzerland, 35–49. doi:10.1007/978-3-031-97629-2_3
- [4] Samin Aref, Mahdi Mostajabadeh, and Hriday Chheda. 2023. Heuristic Modularity Maximization Algorithms for Community Detection Rarely Return an Optimal Partition or Anything Similar. In *Proceedings of the 23th International*




Figure 4: NMI scores on graphs generated with the ADC-SBM graph generator. All graphs have $n = 1\,000$ vertices, $k = 4$ clusters, $d = 20$ average degree, and $\sigma = 1$ intra-cluster center variance. The left plot varies d_{out} while keeping $\sigma_c = 3$, and the right plot varies σ_c while keeping $d_{out} = 2$.

- Conference on Computational Science (ICCS 2023). Springer Nature Switzerland, 612–626. doi:10.1007/978-3-031-36027-5_48
- [5] Samin Aref, Mahdi Mostajabadeh, and Hriday Chheda. 2024. Bayan algorithm: Detecting communities in networks through exact and approximate optimization of modularity. *Physical Review E* 110, 4 (2024), 044315. doi:10.1103/PhysRevE.110.044315
- [6] V. Arnau, S. Mars, and I. Marín. 2004. Iterative Cluster Analysis of Protein Interaction Data. *Bioinformatics* 21, 3 (September 2004), 364–378. doi:10.1093/bioinformatics/bti021
- [7] David A. Bader, Henning Meyerhenke, Peter Sanders, and Dorothea Wagner (Eds.). 2013. *Graph Partitioning and Graph Clustering, 10th DIMACS Implementation Challenge Workshop*, Georgia Institute of Technology, Atlanta, GA, USA, February 13–14, 2012. *Proceedings. Contemporary Mathematics*, Vol. 588. American Mathematical Society. doi:10.1090/conm/588
- [8] Aritra Bhowmick, Mert Kosan, Zexi Huang, Ambuj Singh, and Sourav Medya. 2024. DGCLUSTER: A Neural Framework for Attributed Graph Clustering via Modularity Maximization. In *Proceedings of the 38th AAAI Conference on Artificial Intelligence (AAAI 2024)*. AAAI Press, Washington, DC, USA, 11069–11077. doi:10.1609/aaai.v38i10.28983
- [9] Sonja Biedermann, Monika Henzinger, Christian Schulz, and Bernhard Schuster. 2018. Memetic Graph Clustering. In *Proceedings of the 17th International Symposium on Experimental Algorithms (SEA 2018)*. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 3:1–3:15. doi:10.4230/LIPIcs.SEA.2018.3
- [10] Christopher Blöcker, Chester Tan, and Ingo Scholtes. 2024. The map equation goes neural: Mapping network flows with graph neural networks. In *Proceedings of the 38th Conference on Neural Information Processing Systems (NeurIPS 2024)*. Curran Associates, Inc., 17428–17456. doi:10.52202/079017-0554
- [11] Vincent D. Blondel, Jean-Loup Guillaume, Renaud Lambiotte, and Etienne Lefebvre. 2008. Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment* 2008, 10 (2008), P10008. doi:10.1088/1742-5468/2008/10/P10008
- [12] Ulrik Brandes, Daniel Dellinger, Marco Gaertler, Robert Görke, Martin Hoefer, Zoran Nikoloski, and Dorothea Wagner. 2008. On Modularity Clustering. *IEEE Trans. Knowl. Data Eng.* 20, 2 (2008), 172–188. doi:10.1109/TKDE.2007.190689
- [13] Bingjing Cai, Haiying Wang, Huiru Zheng, and Hui Wang. 2011. Detection of functional modules from protein interaction networks with an enhanced random walk based algorithm. *Int. J. Comput. Biol. Drug Des.* 4, 3 (2011), 290–306. doi:10.1504/IJCDD.2011.041416
- [14] Adil Chhabra, Shai Dorian Peretz, and Christian Schulz. 2025. CluSTRE: Streaming Graph Clustering with Multi-Stage Refinement. In *Proceedings of the 23rd International Symposium on Experimental Algorithms (SEA 2025)*. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 11:1–11:20. doi:10.4230/LIPIcs.SEA.2025.11
- [15] Aaron Clauset, Mark E. J. Newman, and Cristopher Moore. 2004. Finding community structure in very large networks. *Physical Review E* 70, 6 (2004), 066111. doi:10.1103/PhysRevE.70.066111
- [16] John Demme and Simha Sethumadhavan. 2012. Approximate graph clustering for program characterization. *ACM Trans. Archit. Code Optim.* 8, 4 (2012), 21:1–21:21. doi:10.1145/2086696.2086700
- [17] Elizabeth D. Dolan and Jorge J. Moré. 2002. Benchmarking optimization software with performance profiles. *Mathematical Programming* 91, 2 (2002), 201–213. doi:10.1007/s101070100263
- [18] Dušan Džamić, Daniel Aloise, and Nenad Mladenović. 2019. Ascent-descent variable neighborhood decomposition search for community detection by modularity maximization. *Annals of Operations Research* 272, 1 (2019). doi:10.1007/s10479-017-2553-9
- [19] Park K. Fung. 2019. InfoFlow: A Distributed Algorithm to Detect Communities According to the Map Equation. *Big Data and Cognitive Computing* 3, 3 (2019), 42. doi:10.3390/bdcc3030042
- [20] Marco Gaertler. 2004. Clustering. In *Network Analysis: Methodological Foundations [outcome of a Dagstuhl seminar, 13–16 April 2004] (Lecture Notes in Computer Science, Vol. 3418)*, Ulrik Brandes and Thomas Erlebach (Eds.). Springer, 178–215. doi:10.1007/978-3-540-31955-9_8
- [21] M. Girvan and M. E. J. Newman. 2002. Community structure in social and biological networks. *Proceedings of the National Academy of Sciences* 99, 12 (2002), 7821–7826. arXiv:https://www.pnas.org/doi/pdf/10.1073/pnas.122653799 doi:10.1073/pnas.122653799
- [22] Gurobi Optimization, LLC. 2026. Gurobi Optimizer Reference Manual. <https://www.gurobi.com>
- [23] Alexandre Hollocou, Julien Maudet, Thomas Bonald, and Marc Lelarge. 2017. A Streaming Algorithm for Graph Clustering. arXiv preprint. doi:10.48550/arXiv.1712.04337
- [24] Di Jin, Zhizhi Yu, Pengfei Jiao, Shirui Pan, Dongxiao He, Jia Wu, Philip S. Yu, and Weixiong Zhang. 2023. A Survey of Community Detection Approaches: From Statistical Modeling to Deep Learning. *IEEE Trans. Knowl. Data Eng.* 35, 2 (2023), 1149–1170. doi:10.1109/TKDE.2021.3104155
- [25] Ravi Kannan, Santosh S. Vempala, and Adrian Vetta. 2004. On clusterings: Good, bad and spectral. *J. ACM* 51, 3 (May 2004), 497–515. doi:10.1145/990308.990313
- [26] J. Li, S. Lai, Z. Shuai, Y. Tan, Y. Jia, M. Yu, Z. Song, X. Peng, Z. Xu, Y. Ni, H. Qiu, J. Yang, Y. Liu, and Y. Lu. 2024. A Comprehensive Review of Community Detection in Graphs. *Neurocomputing* 600 (2024), 128169. doi:10.1016/j.neucom.2024.128169
- [27] Jiakang Li, Songning Lai, Zhihao Shuai, Yuan Tan, Yifan Jia, Mianyang Yu, Zichen Song, Xiaokang Peng, Ziyang Xu, Yongxin Ni, Haifeng Qiu, Jiayu Yang, Yutong Liu, and Yonggang Lu. 2024. A comprehensive review of community detection in graphs. *Neurocomputing* 600 (2024), 128169. doi:10.1016/J.NEUCOM.2024.128169
- [28] Ke Liang, Lingyuan Meng, Meng Liu, Yue Liu, Wenxuan Tu, Siwei Wang, Sihang Zhou, Xinwang Liu, Fuchun Sun, and Kunlun He. 2024. A Survey of Knowledge Graph Reasoning on Graph Types: Static, Dynamic, and Multi-Modal. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 12 (2024), 9456–9478. doi:10.1109/TPAMI.2024.3417451
- [29] Fanzhen Liu, Shan Xue, Jia Wu, Chuan Zhou, Wenbin Hu, Cécile Paris, Surya Nepal, Jian Yang, and Philip S. Yu. 2020. Deep Learning for Community Detection: Progress, Challenges and Opportunities. In *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence, IJCAI 2020*, Christian Bessière (Ed.). ijcai.org, 4981–4987. doi:10.24963/IJCAI.2020/693
- [30] Yunfei Liu, Jintang Li, Yuehe Chen, Ruofan Wu, Erich Wang, Jing Zhou, Sheng Tian, Shuheng Shen, Xing Fu, Changhua Meng, Weiqiang Wang, and Liang Chen. 2024. Revisiting Modularity Maximization for Graph Clustering: A Contrastive Learning Perspective. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD 2024)*. Association for Computing Machinery, New York, NY, United States, 1968–1979. doi:10.1145/3637528.3671967
- [31] Scott McFarling. 1989. Program Optimization for Instruction Caches. In *ASPLOS-III Proceedings - Third International Conference on Architectural Support for*

- Programming Languages and Operating Systems, Boston, Massachusetts, USA, April 3–6, 1989, Joel S. Emer and John L. Hennessy (Eds.). ACM Press, 183–191. doi:10.1145/70082.68200
- [32] Guilherme Henrique Messias, Tiago Gonçalves Botelho, Sylvia Iasulaitis, Murilo Coelho Naldi, and Alan Demétrius Baria Valejo. 2025. A Systematic Review of Graph Neural Network for Node Clustering in Homogeneous Attributed Networks. *IEEE Access* 13 (2025), 215265–215285. doi:10.1109/ACCESS.2025.3646597
- [33] Atefeh Moradan, Andrew Draganov, Davide Mottin, and Ira Assent. 2023. UCoDe: Unified community detection with graph convolutional networks. *Machine Learning* 112, 12 (2023), 5057–5080. doi:10.1007/s10994-023-06402-0
- [34] M. E. J. Newman. 2003. Properties of highly clustered networks. *Phys. Rev. E* 68 (Aug 2003), 026121. Issue 2. doi:10.1103/PhysRevE.68.026121
- [35] M. E. J. Newman. 2006. Modularity and community structure in networks. *Proceedings of the National Academy of Sciences* 103, 23 (2006), 8577–8582. arXiv:https://www.pnas.org/doi/pdf/10.1073/pnas.0601602103 doi:10.1073/pnas.0601602103
- [36] Michael Ovelgönne and Andreas Geyer-Schulz. 2012. An ensemble learning strategy for graph clustering. In *Graph Partitioning and Graph Clustering*. Contemporary Mathematics, Vol. 588. American Mathematical Society, 187–205. doi:10.1090/conm/588/11705
- [37] Leto Peel, Daniel B. Larremore, and Aaron Clauset. 2016. The ground truth about metadata and community detection in networks. *CoRR abs/1608.05878* (2016). arXiv:1608.05878 <http://arxiv.org/abs/1608.05878>
- [38] Pascal Pons and Matthieu Latapy. 2005. Computing Communities in Large Networks Using Random Walks. In *Proceedings of the 20th International Symposium on Computer and Information Sciences (ISCIS 2005)*. Springer Berlin Heidelberg, 284–293. doi:10.1007/11569596_31
- [39] Arnaud Prat-Pérez, David Dominguez-Sal, Josep M. Brunat, and Josep-Lluís Larriba-Pey. 2012. Shaping communities out of triangles. In *Proceedings of the 21st ACM international conference on Information and knowledge management*. Association for Computing Machinery, 1677–1681. doi:10.1145/2396761.2398496
- [40] Arnaud Prat-Pérez, David Dominguez-Sal, and Josep-Lluís Larriba-Pey. 2014. High quality, scalable and parallel community detection for large real graphs. In *Proceedings of the 23rd International Conference on World Wide Web (WWW 2014)*. Association for Computing Machinery, 225–236. doi:10.1145/2566486.2568010
- [41] Martin Rosvall and Carl T. Bergstrom. 2008. Maps of random walks on complex networks reveal community structure. *Proceedings of the National Academy of Sciences* 105, 4 (2008), 1118–1123. doi:10.1073/pnas.0706851105
- [42] Stanislav Sobolevsky and Alexander Belyi. 2022. Graph neural network inspired algorithm for unsupervised network community detection. *Applied Network Science* 7, 1 (2022), 63. doi:10.1007/s41109-022-00500-z
- [43] Stanislav Sobolevsky, Riccardo Campari, Alexander Belyi, and Carlo Ratti. 2014. General optimization technique for high-quality community detection in complex networks. *Physical Review E* 90, 1 (2014), 012811. doi:10.1103/PhysRevE.90.012811
- [44] Vincent A. Traag, Ludo Waltman, and Nees Jan van Eck. 2019. From Louvain to Leiden: guaranteeing well-connected communities. *Scientific Reports* 9, 1 (2019), 5233. doi:10.1038/s41598-019-41695-z
- [45] Anton Tsitsulin, John Palowitch, Bryan Perozzi, and Emmanuel Müller. 2023. Graph Clustering with Graph Neural Networks. *Journal of Machine Learning Research* 24, 127 (2023), 1–21. <http://jmlr.org/papers/v24/20-998.html>
- [46] Anton Tsitsulin, Benedek Rozemberczki, John Palowitch, and Bryan Perozzi. 2022. Synthetic Graph Generation to Benchmark Graph Learning. arXiv preprint. doi:10.48550/arXiv.2204.01376
- [47] Ken Wakita and Toshiyuki Tsurumi. 2007. Finding community structure in mega-scale social networks: [extended abstract]. In *Proceedings of the 16th International Conference on World Wide Web (WWW 2007)*. Association for Computing Machinery, 1275–1276. doi:10.1145/1242572.1242805
- [48] Hong Wang, Yinglong Zhang, Zhangqi Zhao, Zhicong Cai, Xuwen Xia, and Xing Xu. 2025. Simple yet effective heuristic community detection with graph convolution network. *Scientific Reports* 15, 1 (2025), 39249. doi:10.1038/s41598-025-22860-z
- [49] Aasim Ayaz Wani. 2024. Comprehensive analysis of clustering algorithms: exploring limitations and innovative solutions. *PeerJ Comput. Sci.* 10 (2024), e2286. doi:10.7717/PEERJ-CS.2286
- [50] Ying Xu, Victor Olman, and Dong Xu. 2002. Clustering gene expression data using a graph-theoretic approach: an application of minimum spanning trees. *Bioinform.* 18, 4 (2002), 536–545. doi:10.1093/BIOINFORMATICS/18.4.536
- [51] Tian Zhang, Raghu Ramakrishnan, and Miron Livny. 1996. BIRCH: an efficient data clustering method for very large databases. In *Proceedings of the 1996 ACM SIGMOD International Conference on Management of Data* (Montreal, Quebec, Canada) (SIGMOD '96). Association for Computing Machinery, New York, NY, USA, 103–114. doi:10.1145/233269.233324
- [52] Tianqi Zhang, Yun Xiong, Jiawei Zhang, Yao Zhang, Yizhu Jiao, and Yangyong Zhu. 2020. CommDGI: Community Detection Oriented Deep Graph Infomax. In *Proceedings of the 29th ACM International Conference on Information & Knowledge Management (CIKM 2020)*. Association for Computing Machinery, 1843–1852. doi:10.1145/3340531.3412042

Table 3: Graph from the 10th DIMACS implementation challenge. There are no features or ground truth clusters for these graphs.

Instance	<i>n</i>	<i>m</i>
jazz	198	2 742
celegans-metabolic	453	2 025
email	1 133	5 451
PGPgiantcompo	10 680	24 316
adjnoun	112	425
as-22july06	22 963	48 436
astro-ph	16 706	121 251
celegansneural	297	2 148
cond-mat	16 726	47 594
cond-mat-2003	31 163	120 029
cond-mat-2005	40 421	175 691
dolphins	62	159
football	115	613
hep-th	8 361	15 751
karate	34	78
lesmis	77	254
netscience	1 589	2 742
polblogs	1 490	16 715
polbooks	105	441
power	4 941	6 594
cnr-2000	325 557	2 738 969
eu-2005	862 664	16 138 468
in-2004	1 382 908	13 591 473
uk-2002	18 520 486	261 787 258
uk-2007-05	105 896 555	3 301 876 564
road-central	14 081 816	16 933 413
road-usa	23 947 347	28 854 312
chesapeake	39	170
caidaRouterLevel	192 244	609 066
preferentialAttachment	100 000	499 985
smallworld	100 000	499 998
G-n-pin-pout	100 000	501 198

A Instances

B Memory Usage

C Parallel-sequential Comparison

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009